

Code Optimization: Software pipelining

ARCHI 1

Daniela Genius
daniela.genius@lip6.fr

Loop unrolling

Software pipelining

Loop unrolling



```
Loop: lw R4, 0(R5)      #i1
      sll R7, R4, 1     #i2
      sw R7, 0(R5)      #i3
      addiu R5, R5, 4    #i4
      bne R5, R9, Loop  #i5
      nop               #i6
```

loop body

loop management

Unrolling: register renaming


to limit dependencies WAR between two (or more) iterations
to interlace instructions of different iterations



Introduction of Software pipelining

Loop:	lw R4, 0(R5)	#i1	loop body
	sll R7, R4, 1	#i2	
	sw R7, 0(R5)	#i3	
	addiu R5, R5, 4	#i4	loop management
	bne R5, R9, Loop	#i5	
	nop	#i6	

- Different viewpoint :
without renaming, one can
reorder instructions from
several iterations



- Example for all iterations

Introduction of Software pipelining

```

Loop:
  lw R4, 0(R5)      #i1
  sll R7, R4, 1     #i2
  sw R7, 0(R5)      #i3
  addiu R5, R5, 4   #i4
  bne R5, R9, Loop  #i5
  nop               #i6
    
```


Annotations: Red circles highlight R4, R7, R5, and R9. Red arrows show data flow from R4 to R7 and from R5 to R7 and R5.

i1 (0)
i2 (0)
i1 (1)
i3 (0)
i2 (1)
i1 (2)

i3 (1)
i2 (2)
i1 (3)

i3 (2)
i2 (3)
i1 (4)

i3 (3)
i2 (4)
i3 (4)



Introduction of Software pipelining


```

Loop:
lw (R4), 0(R5)      #i1
sll (R7), (R4), 1    #i2
sw (R7), 0(R5)       #i3
addiu (R5), R5, 4    #i4
bne (R5), R9, Loop   #i5
nop                 #i6
    
```

loop body

loop management

other view: without renaming one can reorder instructions from different iterations



$i1(0)$

$i2(0)$

$i1(1)$

$i3(0)$

$i2(1)$

$i1(2)$

$i3(1)$

$i2(2)$

$i1(3)$

$i3(2)$

$i2(3)$

$i1(4)$

$i3(3)$

$i2(4)$

$i3(4)$

different pattern:

$i3(i-2)$

$i2(i-1)$

$i1(i)$


Introduction of Software pipelining

```

Loop:
lw R4, 0(R5)      #i1
sll R7, R4, 1     #i2
sw R7, 0(R5)      #i3
addiu R5, R5, 4   #i4
bne R5, R9, Loop  #i5
nop               #i6
    
```


loop body

loop management



Principle of software pipelining:
build a loop with instructions
stemming from different iterations

other view: without renaming one can reorder
instructions from different iterations



same pattern:

rebuild a loop:


idea: start iteration before
previous iteration has finished

Software pipelining

Same idea as instruction pipeline

- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it.
BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed resources necessary for stage M (registers used) i.e. terminated the subsequent stage $(M+1)$

loop body




sliced loop



Software pipelining

Same idea as instruction pipeline


- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it. BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed resources necessary for stage M (registers used) i.e. terminated the subsequent stage ($M+1$)



Le pipeline logiciel

Same idea as instruction pipeline

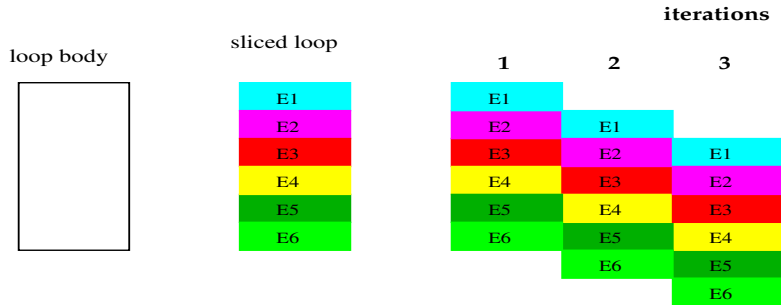
- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it.
BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed ressources necessary for stage M (registers used) i.e. terminated the subsequent stage ($M+1$)



Software pipelining

Same idea as instruction pipeline

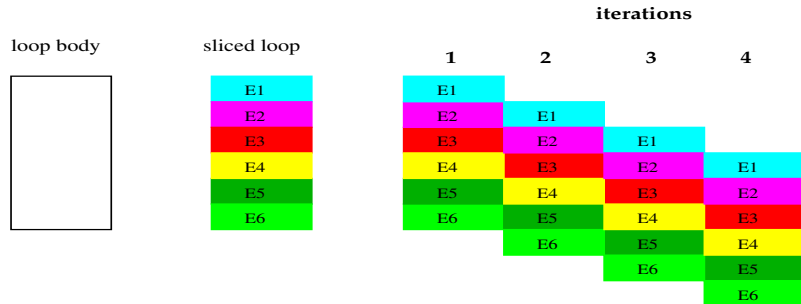
- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it.
BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed resources necessary for stage M (registers used) i.e. terminated the subsequent stage ($M+1$)



Software pipelining

Same idea as instruction pipeline


- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it.
BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed resources necessary for stage M (registers used) i.e. terminated the subsequent stage ($M+1$)



Software pipelining

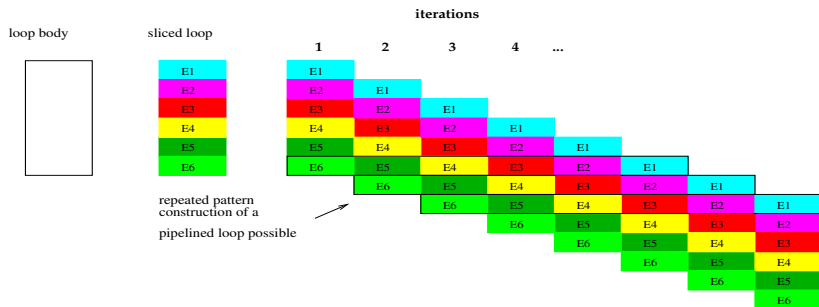
Same idea as instruction pipeline

- ▶ An iteration can be sliced into N stages
- ▶ Stage M of iteration i can be started when iteration $i-1$ has finished it.
BUT program not in // thus stage M of iteration i possible when iteration $i-1$ has freed resources necessary for stage M (registers used) i.e. terminated the subsequent stage ($M+1$)



Software pipelining

- ▶ Pattern that repeats itself, different stage on N subsequent iterations
- ▶ Construction of a pipelined loop becomes possible
- ▶ Part of the initial iteration is executed in the same iteration (cf. in parallel for hardware pipeline)



Software pipelining

- ▶ Loop scheduling a stage from N successive iterations
- ▶ Necessity to add code for filling/emptying the pipeline

sliced loop




pipeline filling code



code folding and
pipelined loop construction



pipeline emptying code



Software pipelining : example

```
# a[i] in R6
# a[N] in R10

loop:
    lw R4, 0(R6) # read elem i
    sll R5, R4, 1 # multiply by 2
    sw R5, 0(R6) # store elem i
    addiu R6, R6, 4
    bne R10, R6, loop
```

Preliminary step : analyze loop to determine dependencies, cycles lost or performance limitation (depending on target architecture)

Step 1 : slicing an iteration into stages

- ▶ Consider loop body and slice into stages
- ▶ Slice where there are stall cycles, dependencies, performance limitations (depending on target architecture)

```
# a[i] dans R6  
# a[N] dans R10  
  
loop:  
    lw R4, 0(R6) # read elem i  
    sll R5, R4, 1 # multiply by 2  
    sw R5, 0(R6) # store elem i  
    addiu R6, R6, 4  
    bne R10, R6, loop
```

Step 1 : slicing an iteration into stages

- ▶ Consider loop body and slice into stages
- ▶ Slice where there are stall cycles, dependencies, Couper là où il y a des cycles de gel, des dépendances, performance limitations (depending on target architecture)
- ▶ Dependencies `lw` → `sll` → `sw`

```
E1 : lw R4, 0(R6)
E2 : sll R5, R4, 1
E3 : sw R5, 0(R6)
```

Step 2 : Build the new loop body

- ▶ Nouveau corps = 1 étape pour n itérations successives si découpage en n étapes
- ▶ Respecter l'ordre des étapes des différentes itérations : itérations les plus anciennes en 1^{er}
- ▶ Si 1^{ère} étape E_1 correspond à itération i : $E_n(i - n + 1), E_{n-1}(i - n + 2), \dots, E_2(i - 1), E_1(i)$
- ▶ $R6 = a[i]$ incremented in each iteration, adjust instructions using it, if required

E3(i-2) : sw R5, ??(R6)

E2(i-1) : sll R5, R4, 1

E1(i) : lw R4, ??(R6)

Step 2 : Build the new loop body

- ▶ R6 incremented by 4 in each iteration, must adjust instructions using it
- ▶ R6 contains address of element i

```
E3(i-2) : sw R5, -8(R6) # elem i-2 is located -2x4  
                                     # bytes in memory before elem i  
E2(i-1) : sll R5, R4, 1  
E1(i)   : lw R4, 0(R6) # elem i, address in R6
```

Step 3 : Insert loop management

- ▶ Stop at adequate moment
- ▶ Stop when last element has been read (stage 1)
- ▶ Stop at address of element N and start at $i=2$
- ▶ R6 must contain address of second element before the loop

```
loop: # de 2 a N-1
    sw R5, -8(R6) # E3(i-2)
    sll R5, R4, 1 # E2(i-1)
    lw R4, 0(R6)  # E1(i)
    addiu R6, R6, 4
    bne R6, R10, loop # R10 contains @tab[N]
    nop
```

Step 3 : insert prologue and epilogue

- ▶ Stop at adequate moment
- ▶ Stop when last element has been read (stage 1)
- ▶ Stop at address of element N and start at $i=2$
- ▶ R6 must contain address of second element before the loop

```
# before : E1(0), E2(0), E1(1)
#         R6 = @tab[2]
loop: # de 2 a N-1
    sw R5, -8(R6) # E3(i-2)
    sll R5, R4, 1 # E2(i-1)
    lw R4, 0(R6)  # E1(i)
    addiu R6, R6, 4
    bne R6, R10, loop # R10 contains @tab[N]
    nop
# after : E3(N-2), E2(N-1), E3(N-1)
```

Reordering (1)

- ▶ Reorder instructions to avoid nop and stall cycles
- ▶ No more intra-iteration RAW dependencies between instructions in loop body
- ▶ Inter-iteration RAW dependencies

```
# before : E1(0), E2(0), E1(1)
#         R6 = @tab[2]'
loop : # i de 2 a N-1
    sw R5, -8(R6) # E3(i-2)
    sll R5, R4, 1 # E2(i-1)
    lw R4, 0(R6)  # E1(i)
    addiu R6, R6, 4
    bne R6, R10, loop # R10 contains @tab[N]
    nop
# after : E3(N-2), E2(N-1), E3(N-1)
```

Reordering (2)

- ▶ Fill the delayed slot : `lw` in delayed slot
- ▶ Eliminate stall cycle `addiu` \rightarrow `bne` : raise the `addiu`
- ▶ 1 iteration = 5 cycles on MIPS32

```
# before : E1(0), E2(0), E1(1)
#           R6 = @tab[2]
loop :
    # i de 2 a N-1
    sw R5, -8(R6)      # E3(i-2)
    addiu R6, R6, 4
    sll R5, R4, 1      # E2(i-1)
    bne R6, R10, loop  # R10 contains @tab[N]
    lw R4, -4(R6)      # E1(i)

# after : E3(N-2), E2(N-1), E3(N-1)
```


Software pipelining : example 2

```
# a[] in R6
# a[N] in R10

loop:
    lw R4, 0(R6) # read val elem i
    sll R5, R4, 1 # val * 2
    add R5, R5, R4 # val * 2 + val == val * 3
    sw R5, 0(R6) # store elem i
    addiu R6, R6, 4
    bne R10, R6, loop
    nop
```

- Analysis of loop body + slicing

Software pipelining : example 2

```
# a[] in R6 '  
# a[N] in R10 '  
  
loop:  
    lw R4, 0(R6) # read val elem i  
    sll R5, R4, 1 # val * 2  
    add R5, R5, R4 # val * 2 + val == val * 3  
    sw R5, 0(R6) # store elem i  
    addiu R6, R6, 4  
    bne R10, R6, loop  
    nop
```

- ▶ Analysis of loop body + slicing
- ▶ RAW dependencies : lw → sll → add → sw
- ▶ Slice into 4 stages

Software pipelining : example 2

- ▶ RAW dependencies : $lw \rightarrow sll \rightarrow add \rightarrow sw$
- ▶ Slice into 4 stages
- ▶ Pipelined loop :

```
# a[i] in R6
# a[N] in R10

loop:
    sw R5, -12(R6) # E4(i-3) store elem i-3
    add R5, R5, R4 # E3(i-2) elem i-2 * 3
    sll R5, R4, 1  # E2(i-1) elem i-1 * (2 + 1)
    lw R4, 0(R6)   # E1(i) read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
    nop
```

Is this loop correct ?

Software pipelining : example 2

```
# a[] in R6
# a[N] in R10
loop:
    sw R5, -12(R6) # E4 : store elem i-3
    add R5, R5, R4 # E3 : elem i-1 * (2 + 1)
    sll R5, R4, 1  # E2 : elem i-2 * 2
    lw R4, 0(R6)   # E1 : read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
```

- ▶ Is this loop correct ?
- ▶ E1 produces R4 used in E2 and E3
- ▶ E2 produces R5 used in E3
- ▶ E3 uses R5 and R4 and produces R5 for E4
- ▶ Explained on blackboard...

Software pipelining : example 2

```
# a[] in R6
# a[N] in R10
loop:
    sw R5, -12(R6) # E4 : store elem i-3
    add R5, R5, R4 # E3 : elem i-1 * (2 + 1)
    sll R5, R4, 1  # E2 : elem i-2 * 2
    lw R4, 0(R6)   # E1 : read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
```

- ▶ Is this loop correct ?
- ▶ E1 produces R4 used in E2 and E3
- ▶ E2 produces R5 used in E3
- ▶ E3 uses R5 and R4 and produces R5 for E4
- ▶ Explained on blackboard...

Software pipelining : example 2

```
# a[] in R6
# a[N] in R10
loop:
    sw R5, -12(R6) # E4 : store elem i-3
    add R5, R5, R4 # E3 : elem i-1 * (2 + 1)
    sll R5, R4, 1  # E2 : elem i-2 * 2
    lw R4, 0(R6)   # E1 : read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
```

- ▶ Is this loop correct ?
- ▶ E1 produces R4 used in E2 and E3
- ▶ E2 produces R5 used in E3
- ▶ E3 uses R5 and R4 and produces R5 for E4
- ▶ Explained on blackboard...

Software pipelining : example 2

```
# a[] in R6
# a[N] in R10
loop:
    sw R5, -12(R6) # E4 : store elem i-3
    add R5, R5, R4 # E3 : elem i-1 * (2 + 1)
    sll R5, R4, 1  # E2 : elem i-2 * 2
    lw R4, 0(R6)   # E1 : read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
```

- ▶ Is this loop correct ?
- ▶ E1 produces R4 used in E2 and E3
- ▶ E2 produces R5 used in E3
- ▶ E3 uses R5 and R4 and produces R5 for E4
- ▶ Explained on blackboard...

Software pipelining : example 2

```
# a[] in R6
# a[N] in R10
loop:
    sw R5, -12(R6) # E4 : store elem i-3
    add R5, R5, R4 # E3 : elem i-1 * (2 + 1)
    sll R5, R4, 1  # E2 : elem i-2 * 2
    lw R4, 0(R6)   # E1 : read elem i

    addiu R6, R6, 4
    bne R10, R6, loop
```

- ▶ Is this loop correct ?
- ▶ E1 produces R4 used in E2 and E3
- ▶ E2 produces R5 used in E3
- ▶ E3 uses R5 and R4 and produces R5 for E4
- ▶ Explained on blackboard...

Software pipelining : back to slicing

- ▶ Each stage must write a value to a register different from the others : if this is not the case, rename
- ▶ Each stage must use the values of the preceding stage and produce only for the subsequent stage : otherwise, insert instructions to pass the values via different registers at each intermediate stage (or revise the slicing into stages)

```
# a[i] in R6, a[N] in R10
loop:
    sw R7, -12(R6)      # E4(i-3) store elem i-3
    add R7, R5, R9      # E3(i-2) elem i-1 * (2 + 1)
    ori R9, R4, 0       # E2(i-1) save R4
    sll R5, R4, 1       # E2(i-1) elem i-2 * 2
    lw R4, 0(R6)        # E1(i) read elem i
    addiu R6, R6, 4
    bne R10, R6, loop
```

Software pipelining : what about internal control flow ?

Example 3



```
# R5 = @tab[i], R7 = @tab[N]
loop:
    lw    R8, 0(R5)
    bgez  R8, endif
    nop
    sub   R8, R0, R8
    sw    R8, 0(R5)
endif:
    addiu R5, R5, 4
    bne   R7, R5, loop
    nop
```

Example 4

```
# R5 = @tab[i], R7 = @tab[N]
loop:
    lw    R8, 0(R5)
    bgez  R8, else
    nop
    addi  R8, R8, -1
    j     suite
    nop
else:
    addi  R8, R8, 1
suite:
    sw    R8, 0(R5)
    addiu R5, R5, 4
    bne   R7, R5, loop
    nop
```



Software pipelining : what about internal control flow ?

- Where to slice the iteration ?



Software pipelining : what about internal control flow ?

- ▶ Cannot separate a branch instruction from the instructions whose execution depends on it, cannot tell the subsequent iteration whether to execute them or not



Summary : Where to slice the iteration ?

- ▶ Consider only the loop body
- ▶ Instructions depending on a branch (control dependency) must be in the same stage as the branch statement
Only dominant blocks are sliced
- ▶ Slice where there is performance loss, dependency of type RAW or stall cycle (depends on architecture)
- ▶ Once slicing realized, verify :
 - ▶ Each stage provides values **exclusively** to the subsequent stage, otherwise add instructions to pass the values or revise the slicing into stages.
 - ▶ Each stage writes to registers different from the other stages, otherwise loss/incoherency of values.