

Multi-Processor Systems

Pirouz Bazargan Sabet

Sorbonne Université - LIP6



Pirouz Bazargan Sabet

Pirouz.Bazargan-Sabet@lip6.fr

February 2021

Introduction

- A Risc can execute on instruction at a cycle
- This rhythm imposes a high exigence on the memory sub-system



Pirouz Bazargan Sabet

February 2021

Outline

- Memory subsystem
- Caches
- Coherency



Pirouz Bazargan Sabet

February 2021

Introduction

What are these requirements ?

- Work at the processor's frequency
- Answer to an instruction fetch request and potentially to a data read or write request in one cycle
- Contain the processor's memory space : 4GB per process
- Reasonable cost



Pirouz Bazargan Sabet

February 2021

Introduction

Is there any memory technology with such characteristics ?

- CMOS Static Memories
 - Work at processor's frequency
 - One cycle per access
 - **KBs to few MBs**



Pirouz Bazargan Sabet

February 2021

Introduction

Is there any memory technology with such characteristics ?

- Hard Disks
 - **Require several million cycles per access**
 - GBs to few TBs



Pirouz Bazargan Sabet

February 2021

Introduction

Is there any memory technology with such characteristics ?

- CMOS Dynamic Memories
 - Work at processor's frequency
 - **Require several cycles per access**
 - **MBs to few GBs**




Pirouz Bazargan Sabet

February 2021

Introduction

Summary



Pirouz Bazargan Sabet

February 2021

Introduction

P

Static

1 cycle
4 KBs

Dynamic

10 cycles
4 MBs

HD

10⁶ cycles
4 GBs

Mean Nbr cycles per access = $\sum c_i P_i$

$$= 1 \times \frac{4K}{4G} + 10 \times \frac{4M - 4K}{4G} + 10^6 \times \frac{4G - 4M}{4G}$$

$$= 1 \times 10^{-6} + 10 \times (10^{-3} - 10^{-6}) + 10^6 \times (1 - 10^{-3})$$

$$= 999\,000 \text{ cycles}$$

Pirouz Bazargan Sabet

February 2021

Introduction

Predict the @ that the processor will access ?

- The previous calculation was based on the assumption that all the @ have the same probability to be accessed

Spatial Locality

Pirouz Bazargan Sabet

February 2021

Introduction

P

Static

1 cycle
4 KBs

Dynamic

10 cycles
4 MBs

HD

10⁶ cycles
4 GBs

Mean Nbr cycles per access = $\sum c_i P_i$

To improve the probability of the data requested by the processor to be present in fast memories, it is necessary to ...

Anticipate
=
Predict

Pirouz Bazargan Sabet

February 2021

Introduction

Spatial locality

P

Static

1 cycle
few KBs

Dynamic

10 cycles
few MBs

HD

10⁶ cycles
few GBs

The probability for a requested data to be present in a memory level :

Hit Rate

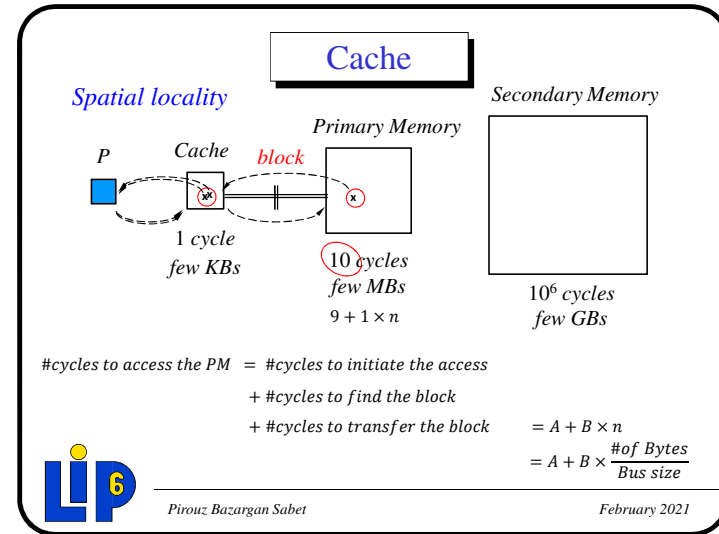
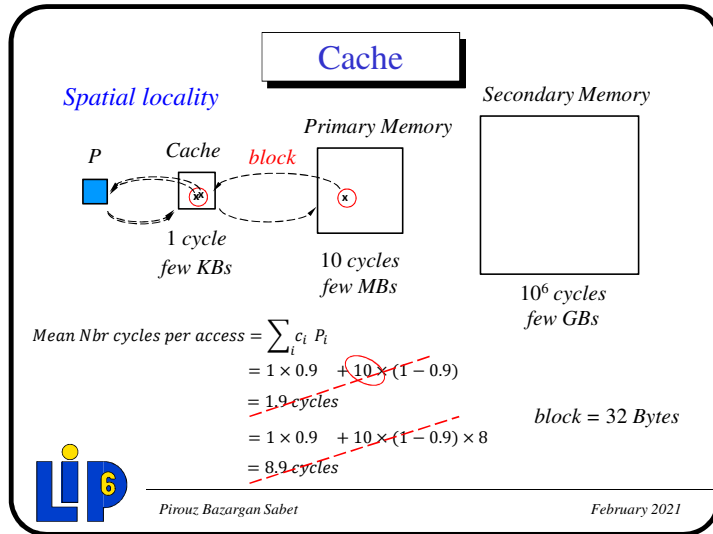
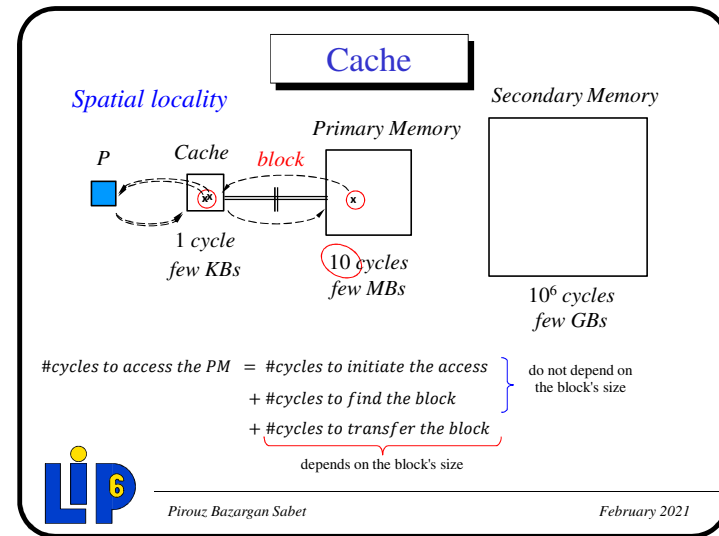
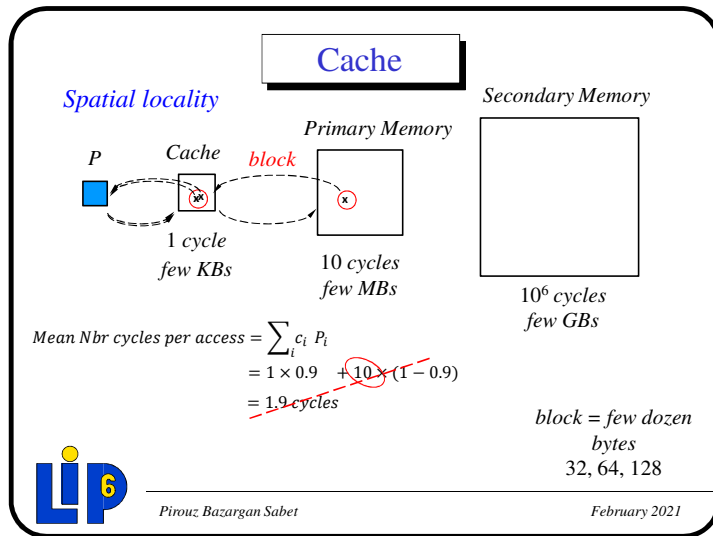
$HitRate_{instruction} = 95\% \text{ à } 98\%$

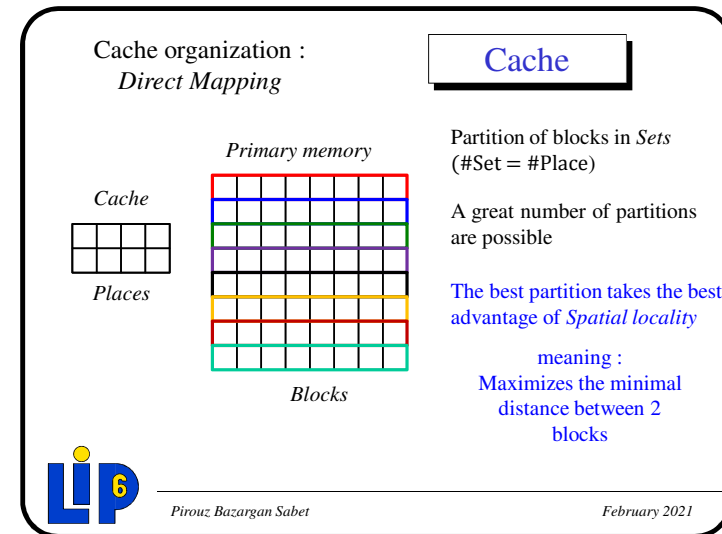
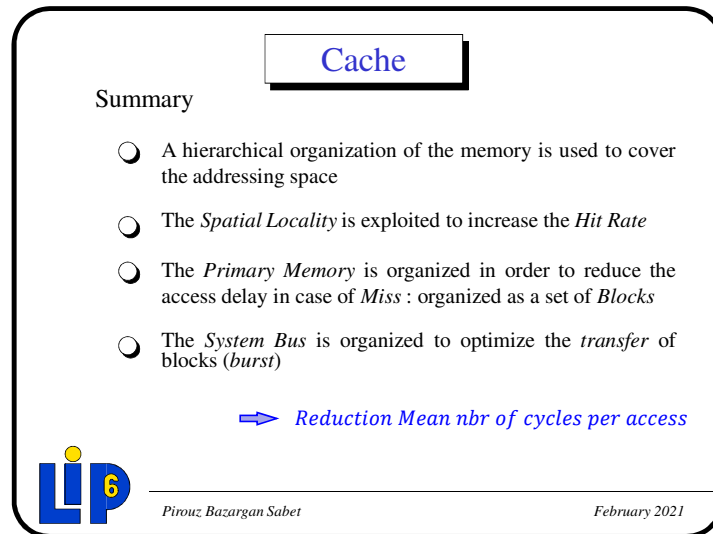
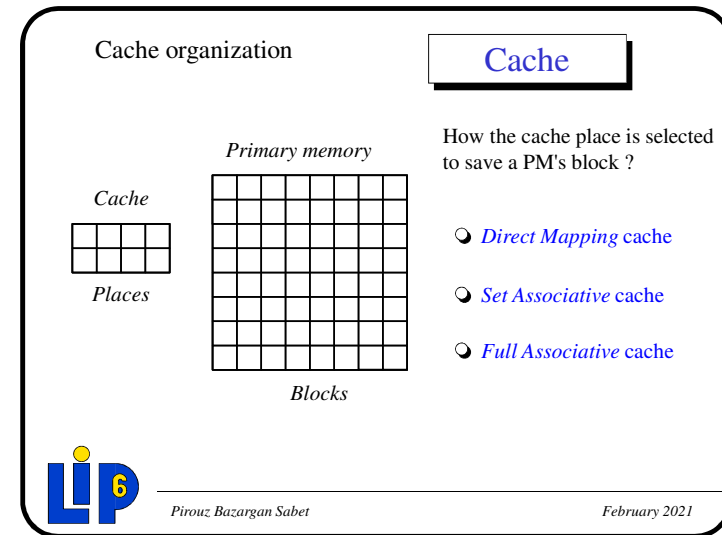
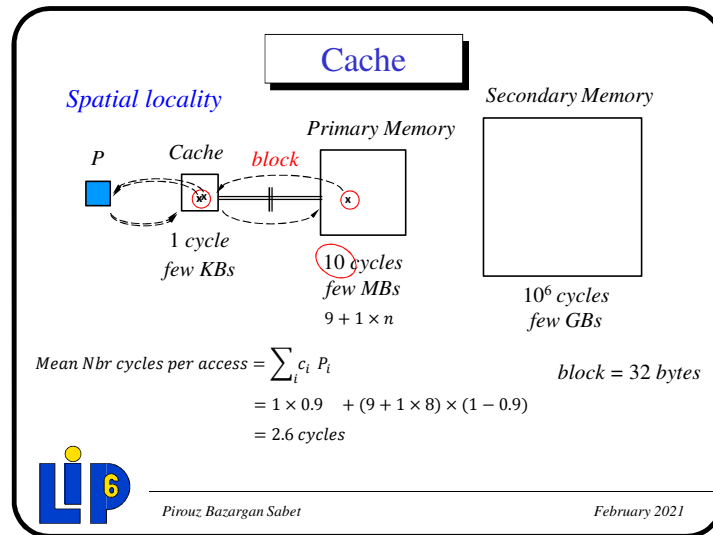
$HitRate_{data} = 90\% \text{ à } 95\%$

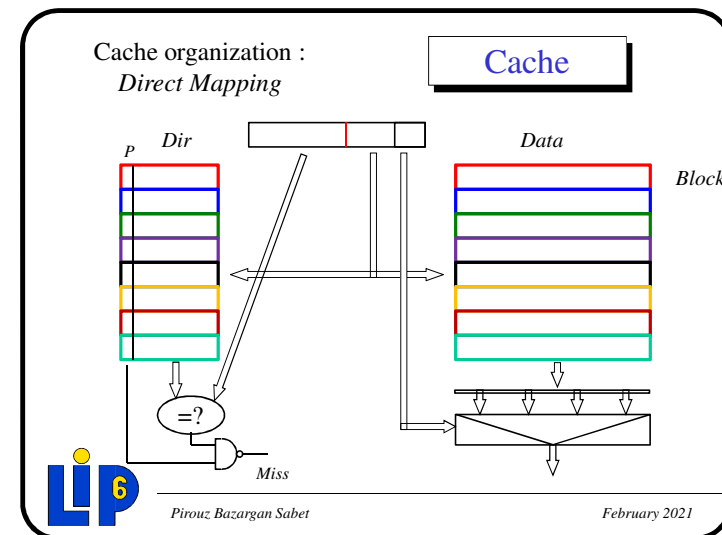
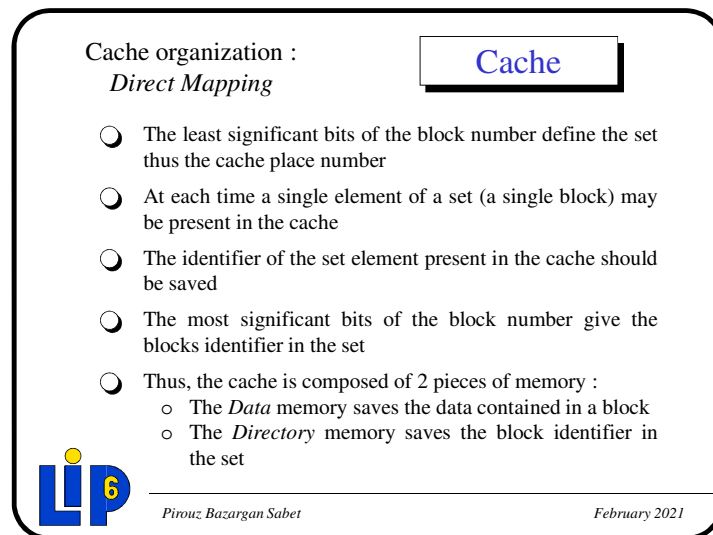
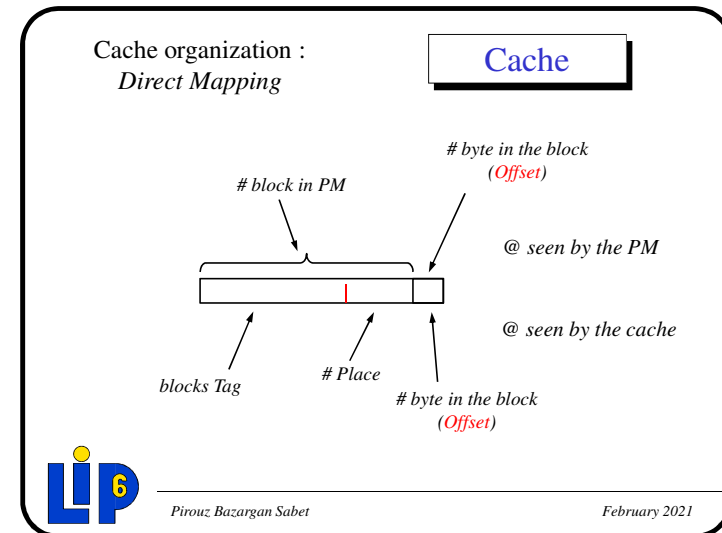
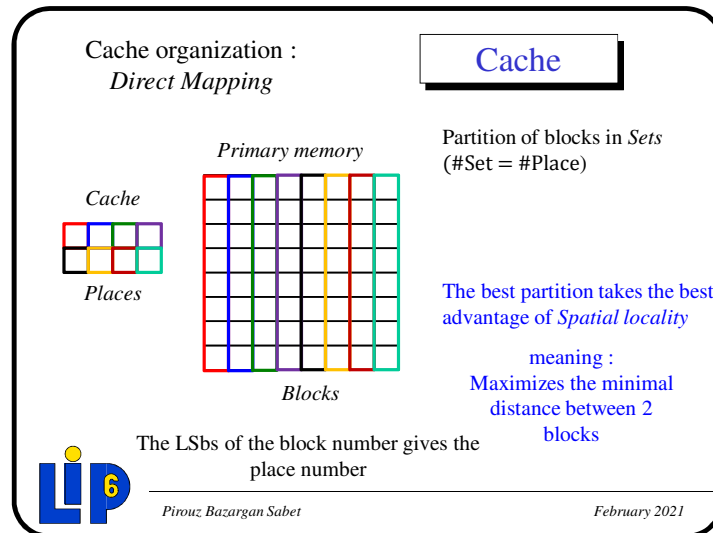
MissRate = 1 - HitRate

Pirouz Bazargan Sabet

February 2021







Cache organization

Cache

How the cache place is selected to save a PM's block ?

- Direct Mapping cache
- Set Associative cache

Pirouz Bazargan Sabet

February 2021

Cache organization :
Set Associatif

Cache

How the cache place is selected to save a PM's block ?

Pirouz Bazargan Sabet

February 2021

Cache organization

Cache

How the cache place is selected to save a PM's block ?

- Direct Mapping cache
- Set Associative cache

Pirouz Bazargan Sabet

February 2021

Cache organization

Cache

How the cache place is selected to save a PM's block ?

- Direct Mapping cache
- Set Associative cache
- Full Associative cache

Pirouz Bazargan Sabet

February 2021

Cache organization :
Set Associatif

Cache

How the cache place is selected to save a PM's block ?

Emplacements

Primary memory

Blocs

Set Associative cache

Replacement strategy : *Pseudo LRU*

LIP 6

Pirouz Bazargan Sabet

February 2021

Cache

Write Strategy

Write Through

P

Cache

Primary memory

LIP 6

Pirouz Bazargan Sabet

February 2021

Cache

Predict the @ that the processor will access ?

Replacement strategy : *LRU*

$P_{@}$

t

Temporal Locality

LIP 6

Pirouz Bazargan Sabet

February 2021

Cache

Write Strategy

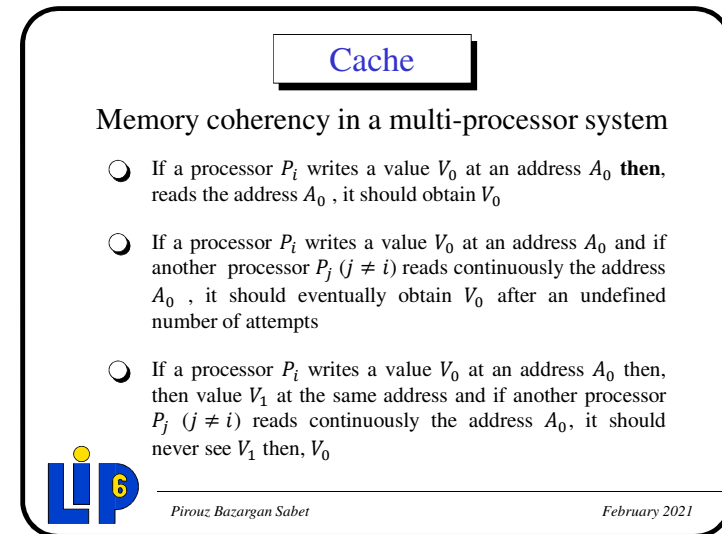
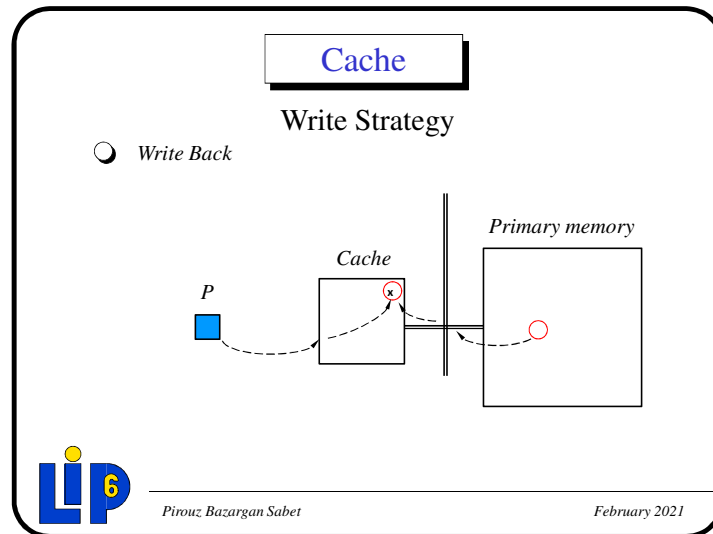
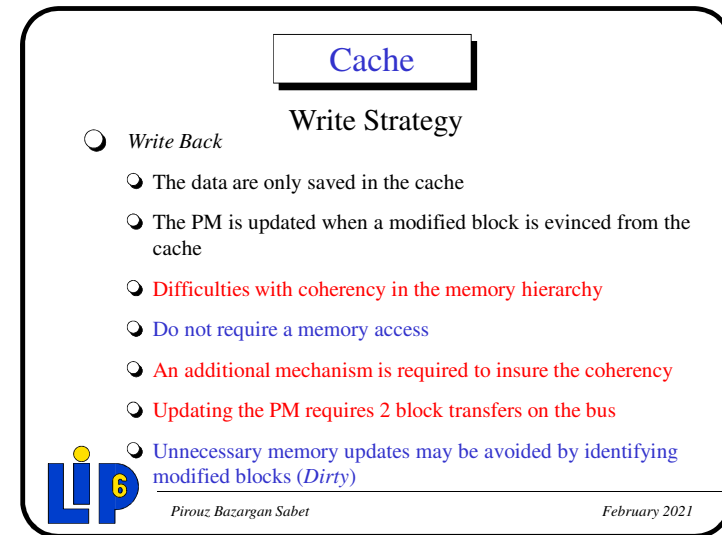
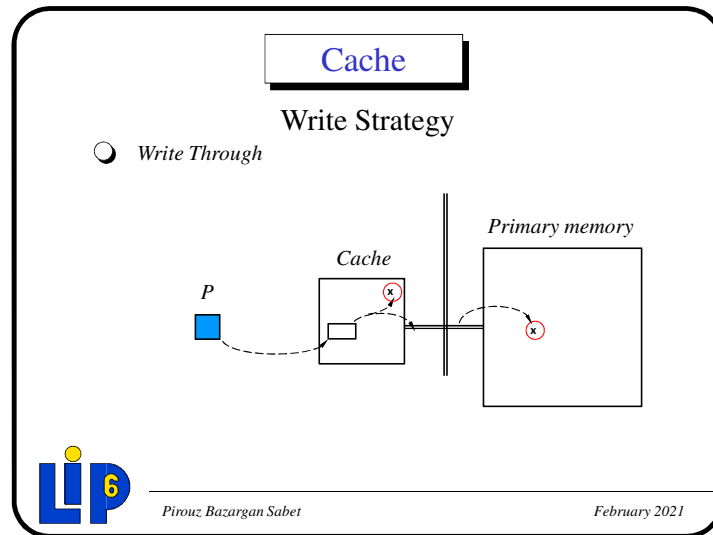
Write Through

- Data are written in the primary memory *through* the cache
- The coherency in the memory hierarchy is fulfilled by construction
- Each single write requires a memory access
- Processor's stall cycles may be avoided by introducing a *Write Buffer*
- The *Write Buffer* contains a small number of places. Stall cycles may be necessary if several writes are received close in time and the *Buffer* goes to be full

LIP 6

Pirouz Bazargan Sabet

February 2021



Cache

How to insure the memory coherency ?

- No cache !!

Some parts of the memory space may be defined as *uncachable*.

Pirouz Bazargan Sabet

February 2021

Pi – Bus

Name	Emitter	Receiver	N bits	Signification
Ck	Ext	CMS	1	Clock : Synchronization signal
ResetN	Ext	CMS	1	Reset : Reset signal – negative logic
Req	M	C	1*M	Request : Request to use the bus
Gnt	C	M	1*M	Grant : Permission to use the bus
A	M	CS	30	Address
Opc	M	S	4	Opcode : Byte selection
Sel	C	S	1*S	Select : Selection of the slave
D	MS	MS	32	Data
Read	M	S	1	Read : Operation type – read or write
Ack	S	MC	3	Acknowledge : Slave's acknowledgment
Lock	M	C	1	Lock : Still need the bus
Tout	C	MS	1	TimeOut : Too long ! – Transaction aborted

Pirouz Bazargan Sabet

February 2021

Cache

How to insure the memory coherency ?

- No cache !!

Some parts of the memory space may be defined as *uncachable*.

- *Uncachable* addresses are known by the cache

or

- An *uncachable* address is notified by the slave during a read

Pirouz Bazargan Sabet

February 2021

Pi – Bus

Ack	S	MC	3	Acknowledge : Slave's acknowledgement
------------	---	----	---	---------------------------------------

RDY	011	Ready
WAT	000	Wait
RTR	100	Retract
UNC		Uncachable

Pirouz Bazargan Sabet

February 2021

Cache

How to insure the memory coherency ?

- No cache !!
Some parts of the memory space may be defined as *uncachable*.
- *Write Through* caches guarantee by construction the memory coherency ... in a mono-processor system and ease the implementation of the coherency in a multi-processor system




Pirouz Bazargan Sabet

February 2021

Cache

Write Strategy

- *Write Back*



A Directory is required for the PM



Pirouz Bazargan Sabet

February 2021

Cache

How to insure the memory coherency ?

- No cache !!
Some parts of the memory space may be defined as *uncachable*.
- *Write Through* caches ease the implementation of memory coherency
- In a *Write Back* cache, when a block is modified the primary memory should be notified that his block is no more up-to-date ... **during the first modification**
From now on, the block is tagged as *Invalid* in the PM



Pirouz Bazargan Sabet

February 2021

Pi – Bus

Op	M	S	4	Opcode : Byte selection
----	---	---	---	-------------------------

WDU	0010	Word Undefined
HW0	1000	Half Word 0
HW1	1010	Half Word 1
BY0	1100	Byte 0
BY1	1101	Byte 1
BY2	1110	Byte 2
BY3	1111	Byte 3
NOP	0000	No Operation
INV		Invalidate



Pirouz Bazargan Sabet

February 2021

Cache

How to insure the memory coherency ?

- No cache !!
Some parts of the memory space may be defined as *uncachable*.
- *Write Through* caches ease the implementation of memory coherency
- In a *Write Back* cache, when a block is modified the primary memory should be notified that his block is no more up-to-date ... *during the first modification*
From now on, the block is tagged as *Invalid* in the PM

An additional mechanism should be implemented in the cache to *monitor* the transfers on the bus. *Snoopy*



Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WT Cache

- The *Snoopy* monitors the *transfers* on the bus
 - *Write* transfers
 - Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
 - Extract the block (read the *Data* – processor's request should wait – be blocked)
 - Modify the block with the data received from the bus
 - Save the modified block in the cache (write into *Data* – processor's request should wait – be blocked)

requires 3 or 4 cycles – *Write* transfer 1 cycle




Pirouz Bazargan Sabet

February 2021

Cache

Cache organization




Pirouz Bazargan Sabet

February 2021

Cache

Cache organization



A *Fifo* can be inserted to save a series (small number) of write transfers close in the time



Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WT Cache

- The Snoopy monitors the *transfers* on the bus
- Write transfers

What if the *Fifo* is near to overflow ?

- Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
- Invalidate the block if present (update the *Directory* – processor's requests should wait – be blocked)



requires 2 cycles – Write transfer 1 cycle

Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WT Cache

- The Snoopy monitors the *transfers* on the bus
- Invalidate transfers

- Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
- Invalidate the block if present (update the *Directory* – processor's requests should wait – be blocked)



requires 2 cycles – Invalidate transfer 1 cycle

Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WT Cache

- The Snoopy monitors the *transfers* on the bus
- Write transfers

What if the *Fifo* is still near to overflow ?

- Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
- Invalidate the cache (update the *Directory* – processor's requests should wait – be blocked)



Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WB Cache

- The Snoopy monitors the *transfers* on the bus
- Write transfers

- Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
- Extract the block (read the *Data* – processor's request should wait – be blocked)
- Modify the block with the data received from the bus
- Save the modified block in the cache (write into *Data* – processor's request should wait – be blocked)



requires 3 or 4 cycles – Write transfer 1 cycle

Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WB Cache

- The Snoopy monitors the *transfers* on the bus
 - Write transfers
 - Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
 - Invalidate the block if present (update the *Directory* – processor's requests should wait – be blocked)



requires 2 cycles – Write transfer 1 cycle

Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WB Cache

- The Snoopy monitors the *transfers* on the bus
 - Invalidate transfers
 - Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
 - Invalidate the block if present (update the *Directory* – processor's requests should wait – be blocked)



requires 2 cycles – Invalidate transfer 1 cycle

Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WB Cache

- The Snoopy monitors the *transfers* on the bus
 - Write transfers
 - What if the *Fifo* comes to is near to overflow ?
 - Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
 - Invalidate the cache (update the *Directory* – processor's requests should wait – be blocked)



Pirouz Bazargan Sabet

February 2021

Cache

Snoopy – WB Cache

- The Snoopy monitors the *transfers* on the bus
 - Aborted transactions (*Retract*)
 - Check if the address of the transfer matches a block present in the cache (check the *Directory* – processor's request should wait – be blocked)
 - Request send to *Bus Interface Controller* to update the memory (*Update Transaction*)




Pirouz Bazargan Sabet

February 2021

Cache

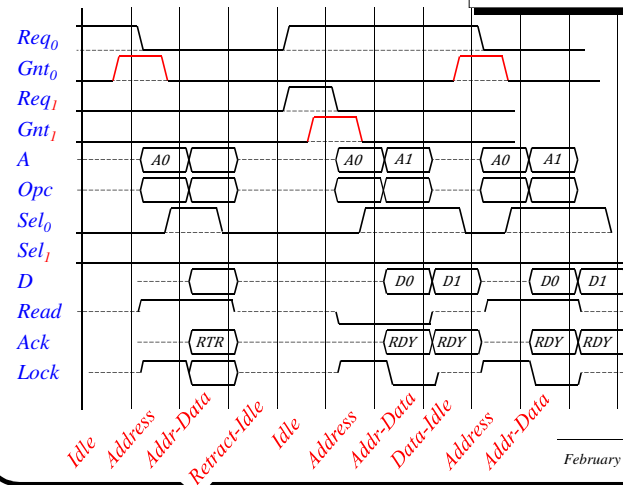
Cache organization : *WB*



Pirouz Bazargan Sabet

February 2021

Pi – Bus : *Retract*



February 2021