

Benchmarking of linear solvers : link between simulation and high performance computing

Anatole Vercelloni and Méline Trochon

Supervisors : Pierre Jolivet and Théo Mary

- 1 Introduction
- 2 The 2D Poisson equation
- 3 MUMPS
- 4 Benchmarks
 - Symmetric permutation
 - Block Low-Rank (BLR)
 - Iterative Refinement
- 5 Conclusion

Introduction




Figure: Solution of a 2D Poisson equation

The 2D Poisson equation

Let us consider the following problem:

$$\begin{cases} \text{find } \varphi \in H^1(\Omega) \text{ s.t.} \\ -\partial_x^2 \varphi - \partial_y^2 \varphi = f \text{ on } \Omega \\ \varphi = g \text{ on } \partial\Omega \end{cases} \quad (1)$$

Thanks to Taylor's expansion, we have :

$$\partial_x^2 \varphi = \frac{\varphi_{i-1,j} - \varphi_{i+1,j} + 2\varphi_{i,j}}{h^2} \quad \text{with } i, j = 1, \dots, n$$

Into a linear system

$$Ax = h^2 f \quad (2)$$

$$\text{with } A = \begin{pmatrix} 4 & -1 & 0 & \cdots \\ -1 & 4 & -1 & \cdots \\ \vdots & \ddots & \ddots & \vdots \\ 0 & \cdots & -1 & 4 \end{pmatrix} \quad \text{and } f = \begin{pmatrix} f_{0,0} & f_{1,0} & f_{2,0} & \cdots \\ f_{0,1} & f_{1,1} & f_{2,1} & \cdots \\ \vdots & \ddots & \ddots & \vdots \\ f_{0,n} & \cdots & f_{n-1,n} & f_{n,n} \end{pmatrix}$$

Sparse linear system

We can remark that A is:

- symmetric (in some case)
- sparse
- low rank properties

- The 2D Poisson equation matrix:
(1002001×1002001) with 7006001 nnz
- The 3D Poisson equation matrix:
(1560896×1560896) with 23091886 nnz

- Single and double precision (Smumps, Dmumps)
- ε -machine
- Scaled residual :

$$r = \frac{\|Ax - b\|_{\infty}}{\|A\|_{\infty} \|x\|_{\infty}} \quad (3)$$

MULTifrontal and Massively Parallel sparse and direct Solver (MUMPS)

- A direct solver for sparse matrices




Figure: The format used in MUMPS

MULTifrontal and Massively Parallel sparse and direct Solver (MUMPS)

- A direct solver for sparse matrices
- Three different phases :
- Analysis phase




Figure: Illustration of a reordering method (AMD) to reduce the fill-in

MUltifrontal and Massively Parallel sparse and direct Solver (MUMPS)

- A direct solver for sparse matrices
- Three different phases :
 - Analysis phase
 - Factorization phase

MUltifrontal and Massively Parallel sparse and direct Solver (MUMPS)

- A direct solver for sparse matrices
- Three different phases :
 - Analysis phase
 - Factorization phase
 - Solve and check phase

Symmetric permutation




Figure: Time then Memory depending on the choice of ICNTL(7) for the 2D Poisson equation

Symmetric permutation




Figure: Time then Memory depending on the choice of ICNTL(7) for the 3D Poisson equation

Block Low-Rank (BLR)






Figure: The number of operation, the memory space then the scaled residual depending on ϵ -BLR for the 2D Poisson problem

Block Low-Rank (BLR)

ϵ	1e-2	1e-4	1e-8
percentage of total over estimated FLOP	3.3	3.8	7.2
percentage of total over estimation entries	17.9	23.3	40.1
scaled residual	5.37e-3	5.53e-4	1.5e-7

Table: The percentage of FLOP and entries and the scaled residual for 3D Poisson equation

Iterative Refinement

repeat

Solve $A\Delta x = r$ using the computed factorization

$x = x + \Delta x$

$r = b - Ax$

The computed backward error ω

until $\omega < \alpha$




Figure: Pseudo-code of the iterative refinement (for a positive parameter)

Iterative refinement





Figure: For 2D Poisson equation, the scaled residual then the time of the solve phase depending on the number of step

Iterative refinement





Figure: For 3D Poisson equation, the scaled residual then the time of the solve phase depending on the number of step

Let μ be the value set by the user, and ε the accuracy they want to reach.

- For 2D : AMD (ICNTL(7) = 0) instead of default
- For 3D : default value for ICNTL(7), and depending on the accuracy :

Let μ be the value set by the user, and ε the accuracy they want to reach.

- For 2D : AMD (ICNTL(7) = 0) instead of default
- For 3D : default value for ICNTL(7), and depending on the accuracy :
- $\varepsilon > \mu$: smumps, BLR (set to ε)

Let μ be the value set by the user, and ε the accuracy they want to reach.

- For 2D : AMD (ICNTL(7) = 0) instead of default
- For 3D : default value for ICNTL(7), and depending on the accuracy :
- $\varepsilon > \mu$: smumps, BLR (set to ε)
- $1e-16 < \varepsilon < \mu$: smumps, BLR (set to μ) + iteration step

Let μ be the value set by the user, and ε the accuracy they want to reach.

- For 2D : AMD (ICNTL(7) = 0) instead of default
- For 3D : default value for ICNTL(7), and depending on the accuracy :
- $\varepsilon > \mu$: smumps, BLR (set to ε)
- $1e-16 < \varepsilon < \mu$: smumps, BLR (set to μ) + iteration step
- $\varepsilon \leq 1e-16$: dmumps, BLR (set to μ) + iteration step

Let μ be the value set by the user, and ε the accuracy they want to reach.

- For 2D : AMD (ICNTL(7) = 0) instead of default
- For 3D : default value for ICNTL(7), and depending on the accuracy :
- $\varepsilon > \mu$: smumps, BLR (set to ε)
- $1\text{e-}16 < \varepsilon < \mu$: smumps, BLR (set to μ) + iteration step
- $\varepsilon \leq 1\text{e-}16$: dmumps, BLR (set to μ) + iteration step
- Thank you for listening !